

Interactive Data Analytics: the New Frontier

Sam Madden

madden@csail.mit.edu

With a cast of many...





bigdata
@CSAIL

BIG



Data



When Do You Have a Big Data Problem?

- Too many bytes (Volume)
- Too high a rate (Velocity)
- Too many sources (Variety)



Real Challenge: Understanding Data

**Required interactivity is
poorly supported by today's
data intensive systems**

What does the data look like?

Show me unusual patterns, events, or outliers?

Where are these anomalies and outliers coming from?

Quickly, as data changes, for arbitrary subsets of the data

Three Interactive Analytics Data Processing Tools We've Built

- **MapD**
 - Interactive data exploration
- **SeeDB**
 - Automatic visualization
- **Scorpion**
 - Understanding “why” in aggregate queries

Can work w/ conventional databases but do better with custom data processing engines


MapD: Interactive Large-Scale Visualization using a GPU Database

MAPD (MASSIVELY PARALLEL DATABASE)
USINGGPUSFORREAL-TIME
VISUALIZATIONOFBIGDATA
AND

w/ Todd Mostak


The Need for Interactive Analytics

- First step in analysis is *browsing*
 - Often visualization
- ad-hoc analytics, with millisecond response times



MapD: GPU Accelerated SQL Database


- ***Key insight:*** GPUs have enough memory that a cluster of them can store substantial amounts of data
- Not an accelerator, but a full blown query processor!
- Massive parallelism enables interactive browsing interfaces
 - 4x GPUs can provide > 1 TB/sec of bandwidth
 - 12 Tflops compute
 - Order of magnitude speedups over CPUs, when data is on GPU
- “Shared nothing” arrangement





147,201,658 tweets from Oct 1, 2012 to Nov 6, 2012



Relative intensity of “tornado” on Twitter (with point overlay) from February 29, 2012 to March 1, 2012









MapD Twitter Demo

y'all



▼

Who

Where

+

?

-


↑

≡ ☁ ⚒ # %

Group by: Word



?



tonight school tomorrow
4sq.com day gonna happy
work night game
bad good love time great
life yeah
hate lmao
people today feel job haha
jobs house girl wanna
sleep morning # Words: 473,616,134

1.0M

800k

600k

400k

200k

2011 April July October 2012 April July October 2013 April July



Follow @dataisland 1.1K likes Show more

M

ANATOMY OF A QUERY


SQL Query: SELECT sender, text FROM tweets WHERE lat < 0.0 ORDER BY DIST(lon, lat, 31.3, 3.0)



Hybrid On-Disk/In-Memory Column Store

Id	lat	lon	date	sender	text
0	31.2	-87.1	10-11	bobama	I lol
1	-41.3	16.4	10-14	smadden	@mit

sender	text
bspears	I sing
smadden	@mit



Row Ids
4
1

Gpu 1

Id	lat	lon	text
0	31.2	-87.1	I lol
4	-17.1	46.3	I sing
8	43.1	-93.7	boston

Gpu 2


Id	lat	lon	text
1	-41.3	16.4	@mit
5	53.1	14.3	haha
9	58.4	2.35	happy

Gpu N

Id	lat	lon	text
3	37.9	-97.8	bieber
7	12.3	11.1	je ne
11	28.4	-81.7	pepsi

Next Steps

- Scale out to many nodes, automate layout algorithms
- Add various advanced analytics (e.g., machine learning algorithms)
- Generalize visualization beyond maps




Three Interactive Analytics Data Processing Tools We've Built

- **MapD**
 - Interactive data exploration
- **SeeDB**
 - Automatic visualization
- **Scorpion**
 - Understanding “why” in aggregate queries

Can work w/ conventional databases but do better with custom engines

SeeDB: Visual Recommendations



w/ Manasi Varak, Aditya Parameswaran, Neoklis Polyzotis

Huge number of possible visualizations

Recommending Visualizations

- **How to find relevant visualizations?**
 - Need a **utility** metric
 - Axes: Data, User Preferences, Aesthetics
- **Goal: interactive recommendations?**
 - Scale to large number of rows
 - Manage curse of dimensionality


SeeDB Visualizations

- $V_i = (d : \text{dimension}, m : \text{measure}, f : \text{aggregate})$
- AGGREGATE + GROUP BY queries

```
SELECT d, f(m) FROM table GROUP BY d  
WHERE selection_predicate
```

Naïvely evaluated through sequential scans of dataset

Result: bar chart



Deviation-based Utility Metric

Find visualizations ($d, f(m)$ sets) such that the difference between the query with the `selection_predicate` and with no predicate is maximized


Recall our query template:

```
SELECT d, f(m) FROM table GROUP BY d  
WHERE selection_predicate
```

Example Recommendation (Census Example)

Bar charts

sex vs. AVG(capital-gain)



*Focus on views that
show high variance vs
underlying data set*

```
SELECT gender, AVG(capital-gain)  
WHERE marital-status = NEVER_MARRIED  
GROUP BY gender
```

Whole data
Set


*Huge
space of
related
problems!!!*

Challenge and Solution Sketch

- **Exponential number of possible visualizations**
 - Can plot any set of attributes against any other set!
- **Solutions:**
 - Several optimizations to *batch* queries together, to explore the search space more efficiently
 - Algorithms to *prune* space of visualizations
 - * **Idea: Quickly discarding those of low utility**
 - * **Evaluate visualizations on a small sample of data**
 - Discard ones that perform poorly, and repeat

Time to Find Top 10 Visualizations

opt NO_OPT SHARING COMB COMB_EARLY



AIR: 10 GB flight dataset
DIAB: 1 GB diabetes patient dataset
BANK: 600 MB banking dataset

SeeDB returns results in < 4 s for all data sets vs > 700s for naïve approach

SeeDB: Visual Recommendations

Data Selector

Dataset **census**

Query Selector

Attribute	Operator	Value
marital	=	Never-

Add Filter

Add Comparison Query

Get Recommendations

Visualization Builder

X axis **sex** Y axis **capital** Y aggregate **Avg**

Plot

Bookmark

sex vs. AVG(capital-gain)

sex	AVG(capital-gain)
Female	~380
Male	~650

Users feel that tool greatly improves their ability to find interesting trends

SeeDB Recommendations

workclass vs. AVG(capital-loss)

workclass	AVG(capital-loss)
Adm&Mng	~120
Exe&Prof	~110
Manuf&Tech	~80
Private	~70
Self-employed	~60
Without-pay	~50

education vs. COUNT(*)

education	COUNT(*)
10th	~100
1st-4th	~100
5th-6th	~100
7th-8th	~100
9th	~100
Assoc-voc	~100
Assoc-acdm	~100
HS-grad	~100
Some-college	~100
Bachelors	~100
Masters	~100
Doctorate	~100

education vs. AVG(capital-loss)

education	AVG(capital-loss)
10th	~50
1st-4th	~50
5th-6th	~50
7th-8th	~50
9th	~50
Assoc-voc	~150
Assoc-acdm	~150
HS-grad	~150
Some-college	~150
Bachelors	~150
Masters	~150
Doctorate	~150

sex vs. COUNT(*)

sex	COUNT(*)
Female	~4,000
Male	~16,000

race vs. COUNT(*)

race	COUNT(*)
Am-Indian-Eskimo	~100
Asian-Pac-Islander	~70,000
Black	~50,000
Other	~45,000
White	~55,000

Three Interactive Analytics Data Processing Tools We've Built

- MapD
 - Interactive data exploration
- SeeDB
 - Automatic visualization
- Scorpion
 - Understanding “why” in aggregate queries

Can work w/ conventional databases but do better with custom engines

Scorpion

- After Se sting, now what?




Eugene Wu

- Common
- Need: a

xist

Definition of Why


Given an outlier group, find a *predicate* over the inputs that makes the output no longer an outlier.




Definition of Why

Given an outlier group, find a *predicate* over the inputs that makes the output no longer an outlier.

i = Input Data



Output Visualization




p = predicate


Definition of Why

Given an outlier group, find a *predicate* over the inputs that makes the output no longer an outlier.

i = Input Data



Output Visualization






Removing the predicate makes US no longer an outlier

What are common properties of those records? {Warren Buffet, Tim Cook}

p : Job = CEO

Scorpion Demo

DBWipes + Scorpion! [toggle scorpion](#)



Existing Data Intensive Systems are a Poor Fit For Interactive Analytic Applications

Example: relational databases

- **Not optimized for interactivity**
 - a query that runs in a few seconds
 - disk-optimized (big pages, buffer pool)
 - ok for the optimizer to work for hours
 - synchronous APIs (JDBC)
- **Designed to perform well on a known workload**
 - careful physical tuning
- **Designed to be the “system of record”**
 - → approximation bad!
- **(Traditional) focus on point lookups & transactional updates**
 - these hurt scan (analytic) performance
- ***Similar statements can be made about Hadoop & Spark***

All 3 examples employ custom data processing layers to circumvent these issues

Four Research Opportunities

- 1. Move away from “index first” and up-front load**

Move away from index first

- **TPC-H Scale 10 Load Times on Postgres**
(~10 GB data, on 4 core MacBook Pro w/ SSD)
 - Load: 7 mins (23 MB/sec)
 - Creating keys: 13 mins
 - Indexing: 18 mins

Untenable if just doing a first pass on the data

Opportunity: index & partition data on the fly

(Some work on avoiding loading too – see, e.g., Alagiannis et al, “NoDB”, SIGMOD 2012)

Example: Database cracking

*Index attributes as they
are accessed, instead of
up front!*

Column A

13
16
4
9
2
12
7
1
19
3
14
11
8
6


Idreos et al. “Database Cracking”, CIDR 2007.

Example: Adaptive Partitioning

- **Data partitioning is key to good performance in modern parallel data systems**
 - Read just the partitions you need
 - Partition is expensive (requires shuffling data)
- **Challenge: how to partition?**
 - Typical choice: partition on frequently queried attributes
 - What if those aren't known?
- **Idea: adaptively partition data as it is queried**

w/ Alekh Jindal, Qui Nguyen, Anil Shanbhag , Aaron Elmore, Divy Agarawal, Jorge Quiane Ruiz

Example: Adaptive Partitioning



Whenever a query arrives, choose whether we should re-partition a block or not

Four Research Opportunities

1. Move away from “index first”
2. Build *analytic* engines for main memory


Analytic engines for main memory: Voodoo Parallel IL

with Oscar Moll, Holger Pirk, Yunming Zhang, Saman Amarasinghe, Matei Zaharia

- To optimize across libraries automatically, need to express them in a common intermediate language
- Design a data-parallel IL that:
 - Captures common data processing tasks
 - Allows rich transformations at the level of the IL
 - Maps efficiently to hardware (clusters, CPU, GPU)
- Focus on main memory & interactive performance

Related: Hyper (TU Munich)

The Goal



Example Transformations: Fusing

```
// library function
def scoreFit(data: vec[vec[float]], param: vec[float]) = {
    sum = [0, 0]
    for (d <- data) { sum += dot(d, param)**2 }
}
```

```
// user code
params = [[1, 1], [3, 2]]
for (p <- params) { scoreFit(data, param) }
```



```
for (p <- params) {
    sum = [0, 0]
    for (d <- data) {
        sum += dot(d, p) ** 2
    }
}
```



```
sums = [[0, 0], [0, 0]]
for (d <- data) {
    for ((p, i) <- params) {
        sum[i] += dot(d, p) ** 2
    }
}
```

Example Transformations: Data Representation

```
// select sum(salary) from users where state == "MA"
def query(users: vec[{name:str, salary:int, state:str}]) = {
    sum = 0
    for (u <- users) {
        if (u.state == "MA") { sum += u.salary }
    }
}
```



```
// column-oriented execution
def query(name: vec[str], salary: vec[int], state: vec[str]) = {
    sum = 0
    for (i <- 0..len(users)) {
        if (state[i] == "MA") { sum += salary[i] }
    }
}
```


Voodoo Backend

- Generates parallel code for a variety of hardware
- Takes as input an intermediate representation of *vectors*

```
// column-oriented execution
def query(name: vec[str], salary: vec[int], state: vec[str]) = {
    sum = 0
    for (i <- 0..len(users)) {
        if (state[i] == "MA") { sum += salary[i] }
    }
}
```

- Hardware abstracted by vector size and number of parallel units
 - Works for GPU, multicore, manycore
- Currently acts as a drop-in backend for MonetDB

Voodoo Performance



TPC-H Query 6, Scale Factor 10 (6 GB scan)

Four Research Opportunities

1. Move away from “index first”
2. Build *analytic* engines for main memory
3. **Treat approximation as a first class citizen**
 - Exploit visual properties

Approximate Data Systems


- By operating on *samples* of data, can get big speedups
 - Example: BlinkDB

BlinkDB

Exploratory analytics workload

42 queries, each of aggregates, groups, and filters on a different subset of attributes.

Runtime Vs. Dataset Size



Approximate Data Systems


- By operating on *samples* of data, can get big speedups
 - Example: BlinkDB
- Historically mildly popular database research topic
 - AQUA, Control, SQL Server
 - Never seen much uptake
 - * DB users aren't comfortable with “close enough”
- Challenges:
 - Queries over rare subgroups
 - * **BlinkDB stratifies on popular attributes**
 - How to compute and maintain random samples
 - What type of sampling to use?

Visually-aware sampling

Correct ordering property


Original

$$\mu_i < \mu_j$$



Sample

$$\nu_i < \nu_j$$



Algorithm sketch:
sample groups
whose confidence
intervals overlap;
don't sample
others

$$\mu_1 < \mu_2 < \dots < \mu_k \rightarrow \nu_1 < \nu_2 < \dots < \nu_k$$


Kim et al, *Rapid Sampling For Visualizations with Order Guarantees* VLDB 2015

MIT COMPUTER SCIENCE AND ARTIFICIAL INTELLIGENCE LABORATORY


Visual Sampling Performance on Flight Dataset

Average Delay by Airline

5-6x speedup over a BlinkDB-like system



Visual sampling: don't paint every pixel



Four Research Opportunities

1. Move away from “index first”
2. Build *analytic* engines for main memory
3. Treat approximation as a first class citizen
 - Exploit visual properties
4. **Develop new asynchronous interfaces**
 - “Linked views”, where V1 updates when V2 changes
 - Incremental refresh of visualizations

* Ex. Meteor Framework: “optimistic UI”



Conclusion

Interactive analytics is a new frontier

Huge performance gulf between current data processing systems (cloud-based or otherwise) and what is required, even on simple tasks

As demand for complex analytics and automated inferences/insight grows, **this gap will get worse**

This creates research opportunities

- In memory engines
- Visually aware approximate processing
- Load less, query more
- New interface abstractions